

Background.

The objective of this project is to simulate target detection in a noisy environment. Everything starts with the signal, $S_n, n \geq 1$, which we assume to be known and deterministic. To model the environment, we assume an auto-regressive noise $V_n, n \geq 1$, of order p . More specifically

$$V_n = \sum_{i=1}^p \varrho_i V_{n-i} + W_n, \quad n \geq 1,$$

where $V_i = 0$ for $i \leq 0$; $\varrho_1, \dots, \varrho_p$ are known with $\varrho_1^2 + \dots + \varrho_p^2 < 1$; and $W_n \stackrel{\text{iid}}{\sim} \mathcal{N}(0, \sigma^2)$ with σ^2 known.

The observed process $X_n, n \geq 1$, is defined by

$$X_n = \mu S_n + V_n,$$

where μ is unknown, and determines the signal “strength”.

Let $p_n = \min\{p, n-1\}$, and for $n \geq 1$ define

$$Y_n = X_n - \sum_{i=1}^{p_n} \varrho_i X_{n-i}, \quad R_n = S_n - \sum_{i=1}^{p_n} \varrho_i S_{n-i}.$$

It is not difficult to see that

$$Y_n = \mu R_n + W_n, \tag{1}$$

and $Y_n \stackrel{\text{indep.}}{\sim} \mathcal{N}(\mu R_n, \sigma^2)$. We will refer to it as the “adjusted” signal in this project. Note that all of the procedures we are interested in (defined below) will only depend on Y_n , not X_n , so technically one does not need to generate the underlying process (X_n) at all.

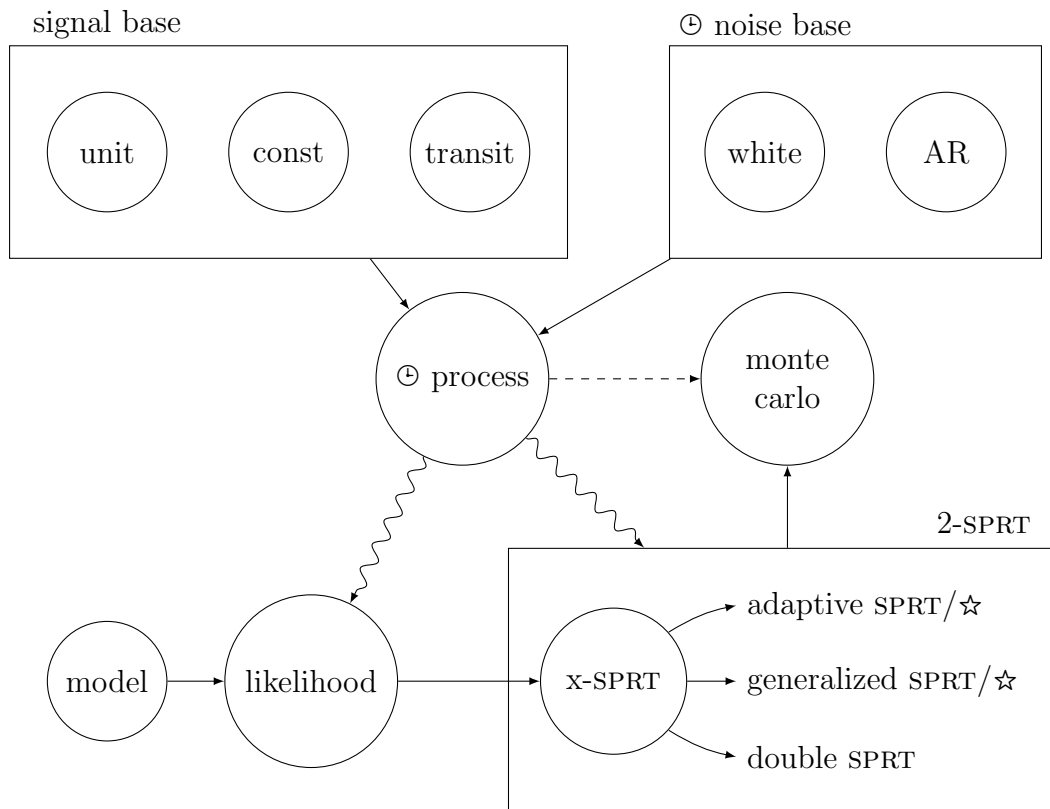
We are now in position to formulate the hypothesis testing problem. Let $\mu_0 = 0$, and $\mu_1 > \mu_0$ be given. We are interested in testing

$$\mathcal{H}_0 : \mu = \mu_0 \quad \text{vs.} \quad \mathcal{H}_1 : \mu \geq \mu_1. \tag{2}$$

In what follows we will also need to introduce auxiliary hypotheses $\{\mathcal{H}_\theta : \mu = \theta\}$.

Project hypotheses.

This project relies heavily on [CRTP](#) (curiously recurring template pattern) structure that serves as a compile-time interface/abstract class/virtual class. The singleton structure was taken from [this post](#) on [stackoverflow](#).



Fundamental structures.

- Signals: `unit_signal`, `constant_signal`, `transitory_signal`.
- Noises: `white_noise`, `auto_regressive_noise`.
- Processes: `process`.
- Hypotheses models: `model`.
- Decision rules: `adaptive_sprt(_star)`, `generalized_sprt(_star)`, `double_sprt`.

CRTP structures.

The following are the base abstract CRTP structures:

- `timed`: for timed random sequences, denoted with \ominus on the diagram.
- `signal_base`: for signals.
- `noise_base`: for noises.
- `observer`: for observers of `process`, denoted with \rightsquigarrow on the diagram.
- `two_sprt`: for SPRT-based decision rules.

2-sprt structure.

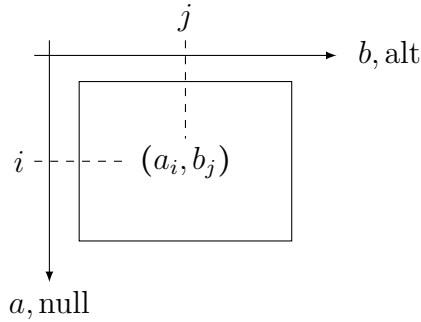
This abstract class represents decision rules (τ, d) , where τ is a stopping time of the form $\tau = \min\{\tau_0, \tau_1\}$, and d is the decision ruling in favor of \mathcal{H}_k if $\tau = \tau_k$, for $k = 0, 1$. The auxiliary stopping times are

$$\tau_0 = \{t \mid T_0 > a\}, \quad \tau_1 = \{t \mid T_1 > b\},$$

where T_0 and T_1 are some scalar-valued SPRT-based statistics, and a, b are thresholds serving as design parameters.

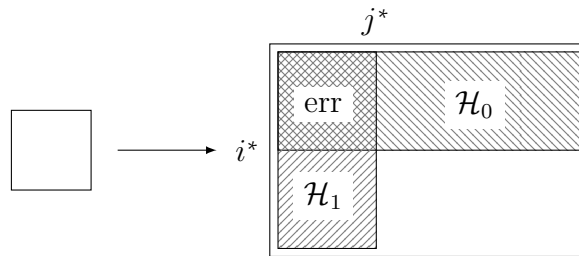
For efficiency reasons, rather than perform simulations for a given pair of thresholds (a, b) , one would construct a grid of thresholds $(a_i, b_j)_{i,j}$ with

$$a_1 < a_2 < \dots < a_m, \quad b_1 < b_2 < \dots < b_n.$$



The simulation then would proceed as long as at least one of the thresholds has not been crossed, or equivalently, as long as (a_m, b_n) has not been crossed. For each pair of thresholds, one would keep track of the observed stopping time, and the decision made.

Let i^* denote the index of the first uncrossed null threshold, a_{i^*} . Let j^* denote the index of the first uncrossed alt threshold, b_{j^*} .



When we start off with the first observation, the entire shaded region in the corresponding matrices will be filled; so each next step only require updating the remaining (rectangular) region of the matrix—for either of the two operating characteristics (OC): stopping time, or decision. So with each next step another Γ -shaped block of the OC matrices will be filled out, yielding an opportunity to optimize the whole process.

Project simulator.

Now to the actual program. The description of simulations to be run are stored in `config`. More specifically, it contains information necessary to create

- `signal`;
- `noise`;
- `monte_carlo`;
- list of `two_sprt`'s;
- list of `run`'s to be performed.

The description of simulation, `run`, contains information about

- `model`;
- list of `simulation_pair`'s;
- list of thresholds for each `two_sprt`.

A `simulation_pair` is the pair of signal strengths: the one simulated in the `process`, denoted here with ν ; and the one analyzed by the `two_sprt`, which we will denote with λ . Regardless of the provided list of `simulation_pair`'s, there are four mandatory pairs to assess basic operating characteristics, summarized in this table:

	$\lambda = \mu_0$	$\lambda = \mu_1$
$\nu = \mu_0$	ESS $_{\mu_0}$	PMS
$\nu = \mu_1$	PFA	ESS $_{\mu_1}$